

Problem Set 5: Advanced Data Structures

Due: Thursday, August 2nd at 10:15 a.m.

Problem 1: Joins on Red-Black Trees (15 points, courtesy of CLR, Problem 14-2, p. 278)

The join operation takes two dynamic sets S_1 and S_2 and an element x such that for any $x_1 \in S_1$ and $x_2 \in S_2$, we have $\text{key}[x_1] \leq \text{key}[x] \leq \text{key}[x_2]$. It returns a set $S = S_1 \cup \{x\} \cup S_2$. Here we'll investigate how to implement this join operation on red-black trees.

- Given a red-black tree T , we store its black-height as the field $\text{bh}[T]$. Argue that this field can be maintained by `RB-Insert` without requiring any extra storage in the tree and without increasing the asymptotic running times. Show that while descending through T , we can determine the black-height of each node we visit in $\Theta(1)$ time per node visited.

We want to implement the operation `RB-Join(T_1, x, T_2)`, which destroys T_1 and T_2 and returns a red-black tree $T = T_1 \cup \{x\} \cup T_2$. Let n be the total number of nodes in T_1 and T_2 .

- Assume without loss of generality that $\text{bh}[T_1] \leq \text{bh}[T_2]$. Describe an $\Theta(\lg n)$ -time algorithm that finds a black node y in T_1 with the largest key among all those nodes whose black-height is $\text{bh}[T_2]$.
- Let T_y be the sub-tree rooted at y . Describe how T_y can be replaced by $T_y \cup \{x\} \cup T_2$ in $\Theta(1)$ time without destroying the binary-search-tree property.
- What color should we make x so that so that red-black properties 1, 2, and 4 are maintained? Describe how property 3 can be restored in $\Theta(\lg n)$ -time.
- Argue that the running time of `RB-Join` is $\Theta(\lg n)$.

Problem 2: Listing all Interval Intersections (10 points, courtesy of Exercise 15.3-4)

Given an interval tree T and an interval i , describe how all intervals in T that overlap i can be listed in $\Theta(\min(n, k \lg n))$, where k is the number of intervals in the output list.

You may temporarily delete entries from the interval tree if necessary, though a snazzier implementation will neither delete them nor mark previously listed nodes in any particular way.

Problem 3: Minimum Gap between Keys (15 points, courtesy of Exercise 15.3-6)

Show how to maintain a dynamic set Q of numbers that supports the operation `Min-Gap`, which gives the magnitude of the difference between the two closest numbers in Q . For example, if $Q = \{1, 5, 9, 15, 18, 22\}$, then `Min-Gap(Q)` would return 3, because

15 and 18 are nearest neighbors. Make the operations `Insert`, `Delete`, `Search`, and `min-gap` as efficient as possible, and analyze their running times. (Hint: Jerry's solution stores three extra fields of information per node, and one of them is called `min-gap[x]`, which tracks the smallest interval between elements in the sub-tree rooted at `x`. To see what other fields are needed, imagine that you need to re-compute the `min-gap` field for the root of the tree, and ask yourself: what information needs to be stored in the root's left and right children in order for me to compute `min-gap[root[T]]`?)